

Program 1 - stats

[Submit Assignment](#)

Due Oct 19 by 11:59pm **Points** 160 **Submitting** a file upload **Available** Sep 24 at 12:01am - Dec 11 at 11:59pm 3 months

Program 1 – CS 344

This assignment asks you to write a bash shell script to compute statistics. The purpose is to get you familiar with the Unix shell, shell programming, Unix utilities, standard input, output, and error, pipelines, process ids, exit values, and signals.

What you're going to submit is your script, called stats.

Overview

NOTE: For this assignment, make sure that you are using Bash as your shell (on Linux, /bin/sh is Bash, but on other Unix O/S, it is not). This is because the Solaris version of Bourne shell has some annoying bugs that are really brought out by this script. Bash can execute any /bin/sh script.


In this assignment you will write a Bourne shell script to calculate averages and medians from an input file of numbers. This is the sort of calculation I might do when figuring out the grades for this course. The input file will have whole number values separated by tabs, and each line of this file will have the same number of values. (For example, each row might be the scores of a student on assignments.) Your script should be able to calculate the average and median across the rows (like I might do to calculate an individual student's course grade) or down the columns (like I might do to find the average score on an assignment).

You will probably need commands like these, so please read up on them: sh, read, expr, cut, head, tail, wc, and sort.

Your script will be called stats. The general format of the stats command is

```
stats {-rows|-cols} [input_file]
```

Note that when things are in curly braces separated by a vertical bar, it means you should choose one of the things; here for example, you must choose either -rows or -cols. The option -rows calculates the average and median across the rows; the option -cols calculates the average and median down the columns. When things are in square braces it means they are optional; you can include them or not, as you choose. If you specify an input_file the data is read from that file; otherwise, it is read from standard input.

Here is a sample run of what your script might return, using an input file called test_file (this particular one can be downloaded [here](#) , note that in Windows, the newline characters may not display as newlines. Move this to your UNIX account, without opening and saving it in Windows, and then cat it out: you'll see the newlines there):

```
% cat test_file
```

```
1      1      1      1      1
9      3      4      5      5
6      7      8      9      7
3      6      8      9      1
3      4      2      1      4
6      4      4      7      7
```

```
% stats -rows test_file
```

```
Average Median
```

```
1      1
5      5
7      7
5      6
3      3
6      6
```

```
% cat test_file | stats -c
```

```
Averages:
```

```
5      4      5      5      4
```

```
Medians:
```

```
6      4      4      7      5
```

```
% echo $?
```

```
0
```

```
% stats
```

```
Usage: stats {-rows|-cols} [file]
```

```
% stats -r test_file nya-nya-nya
```

```
Usage: stats {-rows|-cols} [file]
```

```
% stats -both test_file
```

```
Usage: stats {-rows|-cols} [file]
```

```
% chmod -r test_file
```

```
% stats -columns test_file
```

```
stats: cannot read test_file
```

```
% stats -columns no_such_file
```

```
stats: cannot read no_such_file
```

```
% echo $?
```

```
1
```

Specifications

You must check for the right number and format of arguments to stats. You should allow users to abbreviate -rows and -cols; any word beginning with a lowercase r is taken to be rows and any word beginning with a lowercase c is taken to be cols. So, for example, you would get averages and medians across the rows with -r, -rowwise and -rumplestiltskin, but not -Rows. If the command has too many or too few arguments or if the arguments are of the wrong format you should output an error message to standard error. You should also output an error message to standard error if the input file is specified, but it is not readable.

You should output the statistics to standard output in the format shown above. Be sure all error messages are sent to standard error and the statistics are sent to standard output. If there is any error, the exit value should be 1; if the stats program runs successfully the exit value should be 0.

Your stats program should be able to handle files with any reasonable number of rows or columns. You can assume that each row will be less than 1000 bytes long (because Unix utilities assume that input lines will not be too long), but don't make any assumptions about the number of rows. Think about where in your program the size of the input file matters. You can assume that all rows will have the same number of values; you do not have to do any error checking on this.

You will probably need to use temporary files. For this assignment, the temporary files should be put in the current working directory. (A more standard place for temporary files is in /tmp but don't do that for this assignment; it makes grading easier if they are in the current directory.) *Be sure the temporary file uses the process id as part of its name, so that there will not be conflicts if the stats program is running more than once.* Be sure you remove any temporary files when your stats program is done. You should also use the trap command to catch interrupt, hangup, and terminate signals to remove the temporary files if the stats program is terminated unexpectedly.

All values and results are and must be whole numbers. You may use the expr command to do your calculations, or any other bash shell scripting method, but you may not do the calculations by dropping into another language, like awk, perl, python, or any other language. You may certainly use these other languages for all other parts of the assignment. Note that expr only works with whole numbers. When you calculate the average you should round to the nearest whole number, where half values round up (i.e. 7.5 rounds up to 8). This is the most common form of rounding. You can learn more about rounding methods here (see Half Round Up):

<http://www.mathsisfun.com/numbers/rounding-methods.html> (<http://www.mathsisfun.com/numbers/rounding-methods.html>)

To calculate the median, sort the values and take the middle value. For example, the median of 97, 90, and 83 is 90. The median of 97, 90, 83, and 54 is still 90 - when there are an even number of values, choose the larger of the two middle values.

Your script, stats, must be entirely contained in that file. Do not split this assignment into multiple files or programs.

To make it easy to see how you're doing, you can download the actual grading script here:

[p1gradingscript](#) 

This script is the very one that will be used to assign your script a grade. To compare yours to a perfect solution, you can download here a completely correct run of my stats script that shows what you should get if everything is working correctly:

[p1cleantestscript](#) 

The p1gradingscript itself is a good resource for seeing how some of the more complex shell scripting commands work, too.

Hints

One problem that will be especially challenging is to read in the values from a specified file. The `read` command is exactly what you need (see the man page for `read`). However, `read` is meant to read from standard input and the input file is not necessarily standard input to the `stats` command. You will have to figure out how to get `read` to read from a file. The man page for `sh` has the information you need to figure this out.

Another problem will be calculating the median. There is a straight forward pipelined command that will do this, using `cut`, `sort`, `head`, and `tail`. Maybe you can figure out other ways, too. Experiment!

The `expr` command and the shell can have conflicts over special characters. If you try `expr 5 * (4 + 2)`, the shell will think `*` is a filename wild card and the parentheses mean command grouping.

You have to use backslashes, like this: `expr 5 \'(4 + 2 \')`

Near the top of your program you're going to want to do a conditional test: is the data being passed in as a file, or via `stdin`? One easy way to check for this is to examine the number of parameters used when the script is ran. Once you know, you can store or otherwise process the data correctly, and then pass it onto the calculation parts of your script. In other words, doing the conditional test first, then massaging the data in either form into place in your data structures or temporary files, allows you to write just one version of the calculation, instead of two entirely different blocks of statements!

I HIGHLY recommend that you develop this program directly on the eos-class server. Doing so will prevent you from having problems transferring the program back and forth, which can cause compatibility issues.

If you do see `^M` characters all over your files, try this command:

```
% dos2unix bustedFile
```

Grading

142 points are available from successfully passing the grading script, while the final 18 points will be based on your style, readability, and commenting. Comment well, often, and verbosely: we want to see that you are telling us WHY you are doing things, in addition to telling us WHAT you are doing.